



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

# C/C++ Program Design

## Lab 8, SIMD and OpenMP

于仕琪，廖琪梅



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

# Intel Intrinsics

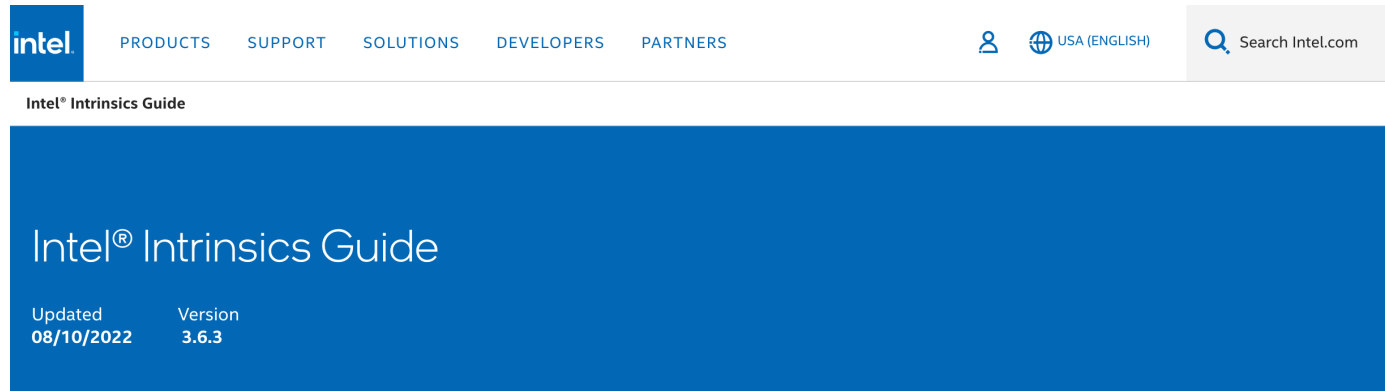
# SIMD@Intel

- MMX: 1997, 8 registers, 64 bits,
- SSE (Streaming SIMD Extensions): 1999, 128 bits
- SSE2: 2000
- SSE3: 2004
- SSSE3: 2006
- SSE4.1: 2006
- SSE4.2
- AVX (Advanced Vector Extensions): 2011, 256 bits
- AVX2: 2013
- AVX-512: 2016



# Intel® Intrinsics Guide

- <https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>



## Instruction Set

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX\_VNNI
- ☐ AVX-512
- ☐ KNC
- ☐ AMX
- ☐ SVML
- ☐ Other

## Categories

- ☐ Application-Targeted
- ☐ Arithmetic

Search Intel Intrinsics

```
void __mm_2intersect_epi32 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2)  vp2intersectd
void __mm256_2intersect_epi32 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2)  vp2intersectd
void __mm512_2intersect_epi32 (__m512i a, __m512i b, __mmask16* k1, __mmask16* k2)  vp2intersectd
void __mm_2intersect_epi64 (__m128i a, __m128i b, __mmask8* k1, __mmask8* k2)  vp2intersectq
void __mm256_2intersect_epi64 (__m256i a, __m256i b, __mmask8* k1, __mmask8* k2)  vp2intersectq
void __mm512_2intersect_epi64 (__m512i a, __m512i b, __mmask8* k1, __mmask8* k2)  vp2intersectq
__m512i __mm512_4dpwssd_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)  vp4dpwssd
__m512i __mm512_mask_4dpwssd_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)  vp4dpwssd
__m512i __mm512_maskz_4dpwssd_epi32 (__mmask16 k, __m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)  vp4dpwssd
__m512i __mm512_4dpwssds_epi32 (__m512i src, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)  vp4dpwssds
__m512i __mm512_mask_4dpwssds_epi32 (__m512i src, __mmask16 k, __m512i a0, __m512i a1, __m512i a2, __m512i a3, __m128i * b)  vp4dpwssds
```

# Load data from memory to registers

```
__m256i _mm256_load_epi32 (void const* mem_addr)
```

vmovdqa32

## Synopsis

```
__m256i _mm256_load_epi32 (void const* mem_addr)
#include <immintrin.h>
Instruction: vmovdqa32 ymm, m256
CPUID Flags: AVX512F + AVX512VL
```

## Description

Load 256-bits (composed of 8 packed 32-bit integers) from memory into `dst.mem_addr` must be aligned on a 32-byte boundary or a general-protection exception may be generated.

## Operation

```
dst[255:0] := MEM[mem_addr+255:mem_addr]
dst[MAX:256] := 0
```

## Latency and Throughput

Architecture	Latency	Throughput (CPI)
Icelake Intel Core	8	0.5
Icelake Xeon	7	0.56
Skylake	8	0.5

```
float * p = ...;
__m256 a;
a = _mm256_load_ps(p);
```

```
__m256i _mm256_load_epi64 (void const* mem_addr)
```

vmovdqa64

```
__m256d _mm256_load_pd (double const * mem_addr)
```

vmovapd

```
__m256h _mm256_load_ph (void const* mem_addr)
```

vmovaps

```
__m256 _mm256_load_ps (float const * mem_addr)
```

vmovaps

```
__m256i _mm256_load_si256 (__m256i const * mem_addr)
```

vmovdqa

# Add operation

```
__m128 _mm_add_ps (__m128 a, __m128 b)  
__m256 _mm256_add_ps (__m256 a, __m256 b)
```

## Synopsis

```
__m256 _mm256_add_ps (__m256 a, __m256 b)  
#include <immintrin.h>  
Instruction: vaddps ymm, ymm, ymm  
CPUID Flags: AVX
```

## Description

Add packed single-precision (32-bit) floating-point elements in `a` and `b`, and store the results in `dst`.

## Operation

```
FOR j := 0 to 7  
    i := j*32  
    dst[i+31:i] := a[i+31:i] + b[i+31:i]  
ENDFOR  
dst[MAX:256] := 0
```

## Latency and Throughput

Architecture	Latency	Throughput (CPI)
Alderlake	2	0.5
Icelake Intel Core	4	0.5
Icelake Xeon	4	0.5
Skylake	4	0.5

**s** is for single precision floating point (float); **d** is for double precision floating point (double)

```
__m256 a, b, c;  
a = _mm256_load_ps(p1 + i);  
b = _mm256_load_ps(p2 + i);  
c = _mm256_add_ps(a, b);
```

**p** is for packed data, all scalars will be in the operation.  
**s** is for scalar, only the first scalar will be involved.

# Store data from registers to memory

```
void _mm_store_pd (double* mem_addr, __m128d a) movq;
```

```
void _mm256_store_pd (double * mem_addr, __m256d a) vmovq;
```

```
void _mm_store_pd1 (double* mem_addr, __m128d a) .
```

```
void _mm_store_ps (float* mem_addr, __m128 a) movaps;
```

## Synopsis

```
void _mm_store_ps (float* mem_addr, __m128 a)
#include <immintrin.h>
Instruction: movaps m128, xmm
CPUID Flags: SSE
```

## Description

Store 128-bits (composed of 4 packed single-precision (32-bit) floating-point elements) from `a` into memory. `mem_addr` must be aligned on a 16-byte boundary or a general-protection exception may be generated.

## Operation

```
MEM[mem_addr+127:mem_addr] := a[127:0]
```

## Latency and Throughput

Architecture	Latency	Throughput (CPI)
Alderlake	1	0.5
Skylake	5	1

```
__m256 c;
float * p = ...;
_mm256_store_ps(p, c);
```

```
void _mm256_store_ps (float * mem_addr, __m256 a) vmovaps;
```

```
void _mm_store_ps1 (float* mem_addr, __m128 a) .
```

```
void _mm_store_sd (double* mem_addr, __m128d a) movq;
```

```
void _mm_store_si128 (__m128i* mem_addr, __m128i a) movd;
```



# ARM Neon Intrinsic



# SIMD@ARM

- Neon: 64 bits and 128 bits
- Helium (or MVE): More instructions
- SVE (Scalable Vector Extension): 128 bits to 2048 bits
- SVE2

# ARM Intrinsics

- <https://developer.arm.com/architectures/instruction-sets/intrinsics/>

**arm**Developer

IP Explorer Documentation Downloads Community Support

Developing on Arm Architecture and Processors Tools and Software

Home / Architectures / Instruction Sets / Intrinsics

## Intrinsics

SIMD ISA

☐ Helium (2471)

☐ Neon (4344)

☐ sve2 (1900)

☐ sve (4140)

Return Base Type

☐ float (2184)

☐ int (4628)

☐ uint (4767)

☐ mve\_pred (177)

☐ void (463)

☐ poly (371)

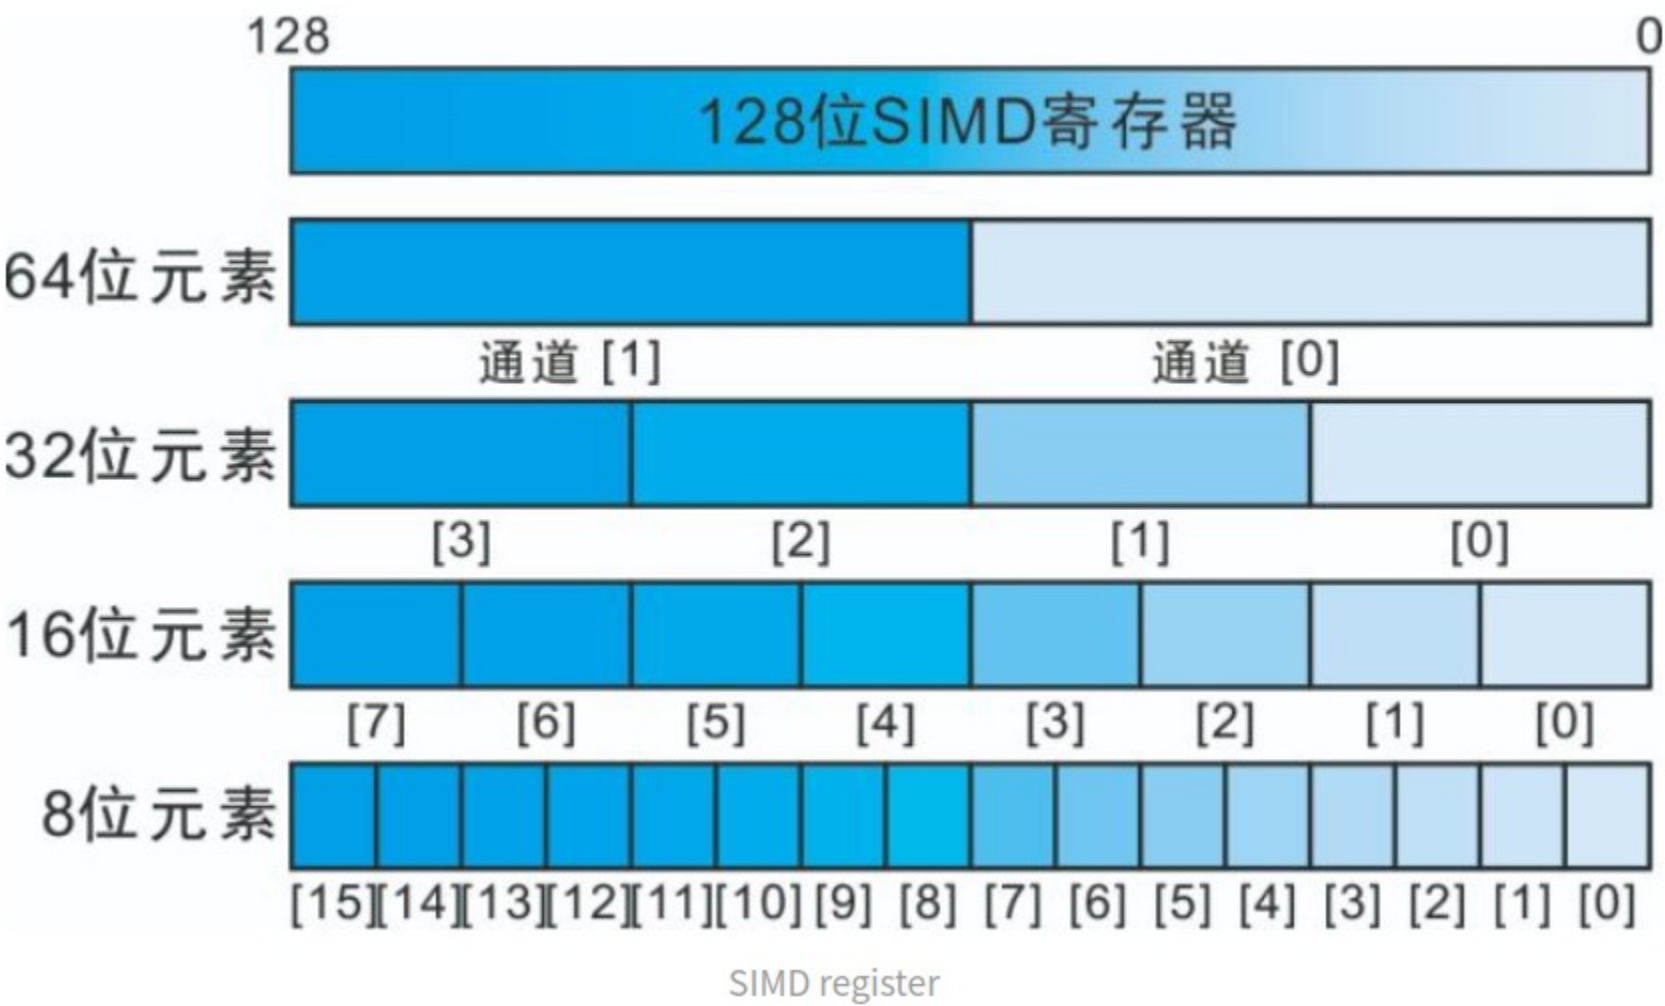
Search All Instructions

### Search Results

Results 1 - 20 of 12855

	SIMD ISA	Return Type	Name	Arguments	Instruction Group
★ ↗	Helium	float16x8_t	[__arm_]vcreateq_f16	(uint64_t a, uint64_t b)	Vector manipulation / Create vector
★ ↗	Helium	float32x4_t	[__arm_]vcreateq_f32	(uint64_t a, uint64_t b)	Vector manipulation / Create vector
★ ↗	Helium	int8x16_t	[__arm_]vcreateq_s8	(uint64_t a, uint64_t b)	Vector manipulation / Create vector
★ ↗	Helium	int16x8_t	[__arm_]vcreateq_s16	(uint64_t a, uint64_t b)	Vector manipulation / Create vector
★ ↗	Helium	int32x4_t	[__arm_]vcreateq_s32	(uint64_t a, uint64_t b)	Vector manipulation /

The NEON Intrinsics can operator on 128-bit registers








<https://manzp.blog.csdn.net/article/details/114686930>















# Load data from memory to registers

	SIMD ISA	Return Type	Name	Arguments	
★ <a href="#">↗</a>	Neon	int8x16_t	vld1q_s8	(int8_t const * ptr)	l
★ <a href="#">↗</a>	Neon	int16x8_t	vld1q_s16	(int16_t const * ptr)	l
★ <a href="#">↗</a>	Neon	int32x4_t	vld1q_s32	(int32_t const * ptr)	l
★ <a href="#">↗</a>	Neon	int64x2_t	vld1q_s64	(int64_t const * ptr)	l
★ <a href="#">↗</a>	Neon	uint8x16_t	vld1q_u8	(uint8_t const * ptr)	l
★ <a href="#">↗</a>	Neon	uint16x8_t	vld1q_u16	(uint16_t const * ptr)	l
★ <a href="#">↗</a>	Neon	uint32x4_t	vld1q_u32	(uint32_t const * ptr)	l
★ <a href="#">↗</a>	Neon	uint64x2_t	vld1q_u64	(uint64_t const * ptr)	l
★ <a href="#">↗</a>	Neon	poly64x2_t	vld1q_p64	(poly64_t const * ptr)	l
★ <a href="#">↗</a>	Neon	float16x8_t	vld1q_f16	(float16_t const * ptr)	l
★ <a href="#">↗</a>	Neon	float32x4_t	vld1q_f32	(float32_t const * ptr)	l
★ <a href="#">↗</a>	Load multiple single-element structures to one, two, three, or four registers. This instruction loads multiple single-element structures from memory and writes the result to one, two, three, or four SIMD&FP registers.				
★ <a href="#">↗</a>					

# Add operation

★ 	Neon	<code>uint8x16_t</code>	<code>vaddq_u8</code>	<code>(uint8x16_t a, uint8x16_t b)</code>	Vector arithmetic / Add / Addition
★ 	Neon	<code>uint16x8_t</code>	<code>vaddq_u16</code>	<code>(uint16x8_t a, uint16x8_t b)</code>	Vector arithmetic / Add / Addition
★ 	Neon	<code>uint32x4_t</code>	<code>vaddq_u32</code>	<code>(uint32x4_t a, uint32x4_t b)</code>	Vector arithmetic / Add / Addition
★ 	Neon	<code>uint64x2_t</code>	<code>vaddq_u64</code>	<code>(uint64x2_t a, uint64x2_t b)</code>	Vector arithmetic / Add / Addition
★ 	Neon	<code>float32x4_t</code>	<code>vaddq_f32</code>	<code>(float32x4_t a, float32x4_t b)</code>	Vector arithmetic / Add / Addition
Description		Floating-point Add (vector). This instruction adds corresponding vector elements in the two source SIMD&FP registers, writes the result into a vector, and writes the vector to the destination SIMD&FP register. All the values in this instruction are floating-point values.			
Results		<code>Vd.4S → result</code>			
This intrinsic compiles to the following instructions:		<code>FADD Vd.4S, Vn.4S, Vm.4S</code>			
Argument Preparation		<code>a → register: Vn.4S</code> <code>b → register: Vm.4S</code>			
Architectures		v7, A32, A64			

# Store data from registers to memory

 	Neon	void	vst1q_u8	(uint8_t * ptr, uint8x16_t val)	Store / Stride
 	Neon	void	vst1q_u16	(uint16_t * ptr, uint16x8_t val)	Store / Stride
 	Neon	void	vst1q_u32	(uint32_t * ptr, uint32x4_t val)	Store / Stride
 	Neon	void	vst1q_u64	(uint64_t * ptr, uint64x2_t val)	Store / Stride
 	Neon	void	vst1q_p64	(poly64_t * ptr, poly64x2_t val)	Store / Stride
 	Neon	void	vst1q_f16	(float16_t * ptr, float16x8_t val)	Store / Stride
 	Neon	void	vst1q_f32	(float32_t * ptr, float32x4_t val)	Store / Stride
Description		Store multiple single-element structures from one, two, three, or four registers. This instruction stores elements to memory from one, two, three, or four SIMD&FP registers, without interleaving. Every element of each register is stored.			
Results		void → result			
This intrinsic compiles to the following instructions:		ST1 {Vt.4S}, [Xn]			
Argument Preparation		ptr → register: Xn val → register: Vt.4S			
Architectures		v7, A32, A64			

Some tips for the example of Week 8

If you compile the source code using “g++ \*.cpp -o main”

```
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ *.cpp -o main
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ ./main
normal: result=9.1, duration = 431ms
unloop: result=9.1, duration = 438ms
NEON is not supported
SIMD: result=0, duration = 0ms
NEON is not supported
SIMD+OpenMP: result=0, duration = 0ms
```

They mean you did not compile the source code with NEON.

If you run the example at Intel CPU, please enable the function call of dotproduct\_avx2() in main.cpp.

```
TIME_START
// result = dotproduct_neon(p1, p2, nSize);
result = dotproduct_avx2(p1,p2,nSize);
TIME_END("SIMD")

TIME_START
// result = dotproduct_neon_omp(p1, p2, nSize);
result = dotproduct_avx2_omp(p1,p2,nSize);
TIME_END("SIMD+OpenMP")
```



2. If you compile it again, the output still mentions AVX2 is not supported.

```
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ -o main *.cpp
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ ./main
normal: result=9.1, duration = 450ms
unloop: result=9.1, duration = 459ms
AVX2 is not supported
SIMD: result=0, duration = 0ms
AVX2 is not supported
SIMD+OpenMP: result=0, duration = 0ms
```

No AVX2 support

The macro WITH\_AVX2 should be enabled when you compile.

```
#ifdef WITH_AVX2
    if(n % 8 != 0)
    {
        std::cerr << "The size n must be a multiple of 8." <<std::endl;
        return 0.0f;
    }

    float sum[8] = {0};
    __m256 a, b;
    __m256 c = _mm256_setzero_ps();

    #pragma omp parallel for
    for (size_t i = 0; i < n; i+=8)
    {
        a = _mm256_load_ps(p1 + i);
        b = _mm256_load_ps(p2 + i);
        c = _mm256_add_ps(c, _mm256_mul_ps(a, b));
    }
    _mm256_store_ps(sum, c);
    return (sum[0]+sum[1]+sum[2]+sum[3]+sum[4]+sum[5]+sum[6]+sum[7]);
#else
    std::cerr << "AVX2 is not supported" << std::endl;
```

To enable the macro by the option -DWITH\_AVX2

```
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ -o main *.cpp -DWITH_AVX2
matoperation.cpp: In function 'float dotproduct_avx2(const float*, const float*, size_t)':
matoperation.cpp:61:34: warning: AVX vector return without AVX enabled changes the ABI [-Wpsabi]
   61 |     __m256 c = _mm256_setzero_ps();
      |                      ^
In file included from /usr/lib/gcc/x86_64-linux-gnu/9/include/emmintrin.h:51,
                  from matoperation.cpp:5:
/usr/lib/gcc/x86_64-linux-gnu/9/include/avxintrin.h:1228:1: error: inlining failed in call to always_inline
 1228 | _mm256_setzero_ps (void)
      | ~~~~~
matoperation.cpp:61:33: note: called from here
   61 |     __m256 c = _mm256_setzero_ps();
      |                  ~~~~~
In file included from /usr/lib/gcc/x86_64-linux-gnu/9/include/emmintrin.h:51,
                  from matoperation.cpp:5:
/usr/lib/gcc/x86_64-linux-gnu/9/include/avxintrin.h:147:1: error: inlining failed in call to always_inline
```

You may get the error message: error:inlining failed .. because you didn't tell the compiler to enable AVX2.

Please use the option -mavx2 to let g++ enable AVX2 support.

```
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ -o main *.cpp -DWITH_AVX2 -mavx2
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ ./main
normal: result=9.1, duration = 447ms
unloop: result=9.1, duration = 448ms
Segmentation fault
```

运行到调用avx2指令时出现段错误

- <https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>

#### Instruction Set

- ☐ MMX
- ☐ SSE
- ☐ SSE2
- ☐ SSE3
- ☐ SSSE3
- ☐ SSE4.1
- ☐ SSE4.2
- ☐ AVX
- ☐ AVX2
- ☐ FMA
- ☐ AVX\_VNNI
- ☐ AVX-512
- ☐ KNC
- ☐ AMX
- ☐ SVM
- ☐ Other

#### Categories

- ☐ Application-Targeted
- ☐ Arithmetic
- ☐ Bit Manipulation
- ☐ Control Flow

Input the name you search

`_mm256_load_ps (float const * mem_addr)`

vmovaps

#### Synopsis

`_mm256_load_ps (float const * mem_addr)`  
`#include <immintrin.h>`  
 Instruction: `vmovaps ymm, m256`  
 CPUID Flags: AVX

It means `_mm256_load_ps()` is in AVX (not AVX2)

#### Description

Load 256-bits (composed of 8 packed single-precision (32-bit) floating-point elements) from memory into `dst.mem_addr` must be aligned on a 32-byte boundary or a general-protection exception may be generated.

#### Operation

`dst[255:0] := MEM[mem_addr+255:mem_addr]`  
`dst[MAX:256] := 0`

#### Latency and Throughput

Architecture	Latency	Throughput (CPI)
Alderlake	7	0.33333333
Icelake Intel Core	7	0.5
Icelake Xeon	7	0.56
Skylake	7	0.5

The option `-mavx` is for AVX, and `-mavx2` is for AVX2

`g++ main.cpp matoperation.cpp -mavx`

`g++ main.cpp matoperation.cpp -mavx2`

You still may get segment fault or wrong results.

- ① For Intel CPU, it's better to use `loadu` & `storeu`, nor `load/store`.
- ② `load` and `store` are for aligned memory only.

```
for (size_t i = 0; i < n; i+=8)
{
    a = _mm256_loadu_ps(p1 + i);
    b = _mm256_loadu_ps(p2 + i);
    c = _mm256_add_ps(c, _mm256_mul_ps(a, b));
}
_mm256_storeu_ps(sum, c);
```

`_loadu` here is for unaligned memory

`_storeu` here is for unaligned memory

Unaligned memory allocation

```
size_t nSize = 200000000;
//float * p1 = new float[nSize](); //the memory is not aligned
//float * p2 = new float[nSize](); //the memory is not aligned

//256bits aligned, C++17 standard
float * p1 = static_cast<float*>(aligned_alloc(256, nSize*sizeof(float)));
float * p2 = static_cast<float*>(aligned_alloc(256, nSize*sizeof(float)));
float result = 0.0f;
```

Aligned memory allocation

### 3. To include different header files by different macros.

```
matoperation.cpp > dotproduct_avx2(const float *, const float *, size_t)
1  #include <iostream>
2  #include "matoperation.hpp"
3
4  #ifdef WITH_AVX2
5  #include <immintrin.h>
6  #endif
7
8
9  #ifdef WITH_NEON
10 #include <arm_neon.h>
11 #endif
12
13 #ifdef _OPENMP
14 #include <omp.h>
15 #endif
```

If you CPU supports AVX2, enable the macro WITH\_AVX2

If you CPU supports AVX2, enable the macro WITH\_NEON

If you you want OpenMP

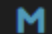
```
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ -o main *.cpp -DWITH_AVX2 -mavx2
maydlee@LAPTOP-U1M00N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ ./main
normal: result=9.1, duration = 456ms
unloop: result=9.1, duration = 449ms
SIMD: result=9.1, duration = 122ms
SIMD+OpenMP: result=9.1, duration = 122ms
```

OpenMP没有启动

You can use `-O3` to gain the maximum speed.

```
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ g++ -o main *.cpp -DWITH_AVX2 -mavx2 -O3
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ ./main
normal: result=9.1, duration = 197ms
unloop: result=9.1, duration = 201ms
SIMD: result=9.1, duration = 33ms
SIMD+OpenMP: result=9.1, duration = 35ms
```

How to do the previous mentioned in CMakeLists.txt

 CMakeLists.txt

```
1  cmake_minimum_required(VERSION 3.12)
2
3  #add_definitions(-DWITH_NEON)
4  add_definitions(-DWITH_AVX2)
5  add_definitions(-mavx)
6  add_definitions(-O3)
7
8  set(CMAKE_CXX_STANDARD 11)
9
10 project(dotp)
11
12 ADD_EXECUTABLE(dotp main.cpp matoperation.cpp)
13
14 find_package(OpenMP)
15 if(OpenMP_CXX_FOUND)
16     message("OpenMP found.")
17     target_link_libraries(dotp PUBLIC OpenMP::OpenMP_CXX)
18 endif()
```

Add some options for the compiler

```
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ mkdir build
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples$ cd build
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples/build$ cmake ..
-- The C compiler identification is GNU 9.4.0
-- The CXX compiler identification is GNU 9.4.0
-- Check for working C compiler: /usr/bin/cc
-- Check for working C compiler: /usr/bin/cc -- works
-- Detecting C compiler ABI info
-- Detecting C compiler ABI info - done
-- Detecting C compile features
-- Detecting C compile features - done
-- Check for working CXX compiler: /usr/bin/c++
-- Check for working CXX compiler: /usr/bin/c++ -- works
-- Detecting CXX compiler ABI info
-- Detecting CXX compiler ABI info - done
-- Detecting CXX compile features
-- Detecting CXX compile features - done
-- Found OpenMP_C: -fopenmp (found version "4.5")
-- Found OpenMP_CXX: -fopenmp (found version "4.5")
-- Found OpenMP: TRUE (found version "4.5")
OpenMP found.
-- Configuring done
-- Generating done
-- Build files have been written to: /mnt/d/mycode/CcodeVS/speedup/examples/build
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples/build$ make
Scanning dependencies of target dotp
[ 33%] Building CXX object CMakeFiles/dotp.dir/main.cpp.o
[ 66%] Building CXX object CMakeFiles/dotp.dir/matoperation.cpp.o
[100%] Linking CXX executable dotp
[100%] Built target dotp
maydlee@LAPTOP-U1MO0N2F:/mnt/d/mycode/CcodeVS/speedup/examples/build$ ./dotp
normal: result=9.1, duration = 270ms
unloop: result=9.1, duration = 280ms
SIMD: result=9.1, duration = 47ms
SIMD+OpenMP: result=0, duration = 11ms
```

You can create a directory for generated files by cmake

To use the file CMakeLists.txt in the parent directory

Make it!

Run it!



# Exercise:

Write a program to add 2 `float` vectors whose size should be more than 1M. You can initialize the two vectors with values like `0.f`, `1.f`, `2.f`, ...

- Use pure C source code and SIMD (AVX2 or NEON) separately, and compare their speeds
- Use OpenMP to speed up the addition. Can you get the correct result?